

```
In [1]: import numpy as np
import pandas as pd
import sklearn

from sklearn import tree
from pandas import *

df = pd.read_csv("movie_metadata.csv")
df = df.dropna()
```

```
In [2]: from sklearn.feature_extraction.text import CountVectorizer
count_vect = CountVectorizer(min_df=1)
```

```
In [3]: from sklearn.preprocessing import LabelEncoder
X = pd.DataFrame()
df = pd.read_csv("movie_metadata.csv")
df = df.dropna()

columnsToEncode = list(df.select_dtypes(include=['category', 'object']))
le = LabelEncoder()
for feature in columnsToEncode:
    try:
        df[feature] = le.fit_transform(df[feature])
    except:
        print('Error encoding ' + feature)
df.head()
```

```
Out[3]:
```


	color	director_name	num_critic_for_reviews	duration	director_facebook_likes	actor
0	1	620	723.0	178.0	0.0	855.0
1	1	538	302.0	169.0	563.0	1000
2	1	1395	602.0	148.0	0.0	161.0
3	1	251	813.0	164.0	22000.0	2300
5	1	62	462.0	132.0	475.0	530.0

5 rows × 28 columns

```
In [4]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

corr = df.select_dtypes(include = ['float64', 'int64']).iloc[:, 1:].corr()
plt.figure(figsize=(15, 15))
sns.heatmap(corr, vmax=1, square=True)
```

Out[4]: <matplotlib.axes._subplots.AxesSubplot at 0x2923357ba90>



```
In [5]: X=df
        y=X['imdb_score']
        #y.apply(np.round)
        X = X.drop(['imdb_score'], axis = 1)

        from sklearn.cross_validation import train_test_split
        from sklearn.preprocessing import StandardScaler

        scaler=StandardScaler()
        X = scaler.fit_transform(X)
        y = np.array(y).astype(int)

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3, random
        _state =190)
```

Linear and Logistic Regression

```
In [28]: from sklearn.linear_model import LogisticRegression
        LRL2 = LogisticRegression(penalty = 'l2')
        LRL2.fit(X_train,y_train)
        L2score = LRL2.score(X_test,y_test)
        print (L2score)
```

0.482697426797

```
In [25]: from sklearn.linear_model import LinearRegression
        model = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=4,
        normalize=True)
        model.fit(X_train,y_train)
        print('Accuracy: ', model.score(X_test, y_test))
```

Accuracy: 0.371803254538

Decision Tree

```
In [18]: #predict and score
        from sklearn import tree
        tree_model = tree.DecisionTreeClassifier(max_depth = 11, min_samples_split=90)
        tree_model.fit(X_train, y_train)
        tree_model.score(X_test, y_test)
```

Out[18]: 0.50044365572315885

SVM

```
In [29]: from sklearn import svm

clf = svm.SVC()
clf.set_params(C=10)
clf.fit(X_train, y_train)
clf.predict(X_test)

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
prediction = clf.predict(X_test)
print(accuracy_score(y_test, prediction))

0.519077196096
```

Linear SVM

```
In [37]: clfLin = svm.SVC(kernel = 'linear')
clfLin.set_params(C=0.5)
clfLin.fit(X_train, y_train)
clfLin.predict(X_test)

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
prediction = clfLin.predict(X_test)
print(accuracy_score(y_test, prediction))

0.512866015972
```

KNN

```
In [34]: from sklearn.neighbors import KNeighborsClassifier

neigh = KNeighborsClassifier(n_neighbors=12)

neigh.fit(X_train, y_train)
neigh.set_params(p=7)
neigh.predict(X_test)
accuracy_score(y_test, neigh.predict(X_test))
```

Out[34]: 0.42236024844720499

Random Forest

```
In [35]: from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(n_estimators=28)
clf = clf.fit(X_train, y_train)
accuracy_score(y_test, clf.predict(X_test))
```

Out[35]: 0.54835847382431235

In []: